

Improving cache allocation for networking applications using fine-grained cache partitioning and unikernels

Jiangqiong Liu
liu3328@purdue.edu
Purdue University
United States

Joshua Repogle
jreprogl@purdue.edu
Purdue University
United States

Vedaant Rajoo
vrajoo@purdue.edu
Purdue University
United States

Sowmya Jayaram Iyer
jayarami@purdue.edu
Purdue University
United States

Keerthana Ashokkumar
kashokku@purdue.edu
Purdue University
United States

Anmol Sahoo
Purdue University
United States
sahoo9@purdue.edu

ABSTRACT

We present a solution for improving the performance of networking applications. The performance of network IO is bound by the delay of accessing main memory and thus processors implement Direct Cache Access (DCA) technology such as Intel DDIO to bypass the main memory and improve performance. But, cache contention between different parts of an application hurts the performance of DDIO and thus we propose a solution that minimizes such interference. We rely on two crucial insights - intra-application cache partitioning to isolate networking and non-networking parts of the code and deploying unikernels to remove interference from the operating system. We present an implementation of this solution using the Intel Resource Director Technology (RDT) for cache partitioning and the Unikraft unikernel library. Finally, we present an evaluation strategy using hardware counter events and report our results over a set of benchmarks.

1 INTRODUCTION

As networking interfaces start breaking the 100Gbps barrier, it is being observed that commodity processors are not able to keep up with processing packets at the rate required to saturate these links. With the end of Moore's law and Dennard scaling, there is an upper limit to how many instructions can be processed in a given interval, which limits the bandwidth achievable on commodity processors. This is further exacerbated by the fact that memory speeds have stayed mostly constant making main memory access the critical bottleneck to processing packets at fast rates. To deal with this problem, modern processors often include features implementing Direct Cache Access (DCA) which allows peripheral devices to directly read and write data from the last-level cache of the processor. Since the peripheral can bypass the entire main memory, this reduces the effect of the slow memory hierarchy on the packet processing critical path and provides a pathway to better performance, at least on paper. But the presence of DCA does not necessarily improve performance and in certain cases may even decrease performance. Typically, since caches are small and there is high contention for cache space, DCA may not be able to place packets efficiently in

the cache. A thorough study of the effects of DCA on networked applications is provided in [2]. While existing solutions to this problem rely on using coarse grained cache partitioning to isolate different applications, they do not consider the cache contention arising from *within* an application. As we know, the computational and networking parts of the same application may have different cache footprints which contend for the same amount of cache space which limits the effectiveness of DCA, since they will evict each others cache lines due to this contention. To tackle this problem, we develop a solution based on two key insights - intra-application cache partitioning and unikernels. We show that these two insights together are necessary to achieve a holistic solution that reduces the effects of cache contention arising from within different parts of the same application. Concretely, our contributions are as follows -

- (1) An interface by which networking applications can demarcate regions of their code that perform network IO
- (2) An implementation that allows executing such applications with fine-grained cache partitioning as a unikernel
- (3) A set of benchmarks and an evaluation strategy for checking the performance of our solution

While our results look promising for some benchmarks (~10% improvement in runtime), they are not consistent across our entire testset. We identify areas of improvement which can improve performance consistently.

2 BACKGROUND AND MOTIVATION

Peripheral devices rely on the use of Direct Memory Access (DMA) to perform IO operations. DMA technology implemented in processors allows peripheral devices to directly write to system main memory through the Peripheral Component Interface-express (PCIe) bus. The typical flow of a network interface card (NIC) peripheral to receive a packet is to write the incoming packet in main memory and then signal the processor via an interrupt about this event. The processor may perform additional processing on the packet, write a new packet to main memory and then signal the NIC to send it out. Given this setup, the latency of accessing main memory is the main bottleneck. At 100Gbps, a processor has 6.72ns to process small packets. But a memory access on the order of 100ns, makes processing packets at line rate prohibitive. To address this problem processors implement Direct Cache Access (DCA) technology. DCA is an umbrella term to refer to technologies that can be used to improve access time of peripheral DMA data, by placing it in the

cache. For example, a simple technique allows a cache prefetcher to prefetch parts of the peripheral memory, thus improving access in predictable scenarios. Intel processors implement a DCA feature named DDIO (Direct-Data IO). In this setup, a NIC can perform DMA directly from/to the cache bypassing the main memory altogether. This prevents the delay incurred by the processor accessing the packet from main memory. Similarly, while writing the packet the NIC can read the packet from the cache, avoiding the latency of the main memory. We describe this process in detail, as it highlights the role that cache contention plays in this situation. When a NIC receives a packet, it writes a cache line (memory address block) into the LLC via PCIe. In this case, DDIO overwrites the cache line if it is present in the cache (PCIe write hit / write update) or allocates a new line if not (PCIe write miss / write allocate). Note that in the first case, DDIO may write the packet anywhere in the cache but in the second case, it may only allocate a packet in a specific region of the cache implying that frequent write allocations will not make effective use of the cache. When the processor signals the NIC to send a cache line containing the packet, the NIC will directly read the packet from the cache using DDIO. Similar to the previous case, if the cache line is present in the cache then it is a hit (PCIe read hit) otherwise it reads the cache line from main memory (PCIe read miss). With this, it becomes possible to see how DDIO may be affected by cache contention. The first problem is called leaky DMA [7], where incoming packets arrive faster than they can be processed thus evicting older packets from memory. While one solution to this is to limit the number of RX descriptors (so that the cache space used by DDIO is bounded) it is only a temporary patch. Small number of RX descriptors increases the packet loss, since fewer packets can be buffered now. Similarly, limiting TX descriptors is also another proposed solution, but reducing the number of TX descriptors leads to inefficient utilization of the PCIe bandwidth. Thus, both of these problems are mitigated by increasing the number of RX/TX descriptors, but this then overloads the cache and leads to reduced DDIO efficiency. This highlights the role that cache contention plays in determining the performance of network IO, since it leads to the eviction of RX/TX descriptors and buffers. A significant source of such cache contention exists within an application itself, since an application maybe performing non-trivial memory bound computation while simultaneously performing network IO. Having motivated the problem, we focus on solving cache-contention between different parts of the application with the intention of improving network performance. We provide a simple motivating example in Fig.1. As can be seen in Fig.1(a), at a certain point in the program's call stack, the cache is occupied by the packets that were received due to network IO. Now, the application executes some computation which displaces these cache lines, as seen in Fig.1(b). Now any future network IO that may need to be performed on these packets, would need to fetch them from main memory again. This clearly illustrates the problem of different parts of the application contending for cache memory and thus affecting networking performance.

3 PROPOSED SOLUTION

To mitigate the problems that we discussed in Section 2, we develop our solution based on two key insights - (1) intra-application

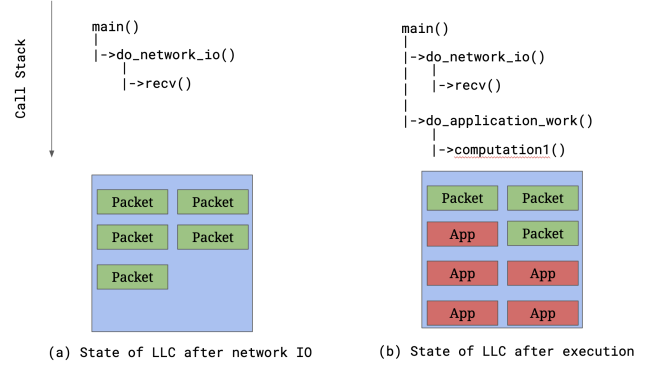


Figure 1: Motivating Example

cache partitioning and (2) using unikernels to remove OS contention. Intra-application cache partitioning refers to performing fine-grained partitioning of the cache, such that different parts of the same application occupy different parts of the cache, thus reducing cache contention. By isolating the parts of the cache in which the RX/TX descriptors and buffers reside, interference from different parts of the application are prevented. Thus, these buffers are not evicted and DDIO performance is not affected. The second idea of using a unikernel is to reduce interference from the Operating System (OS). An Operating System (OS) is responsible for providing an abstraction of the CPU for applications. It maintains multiple in-memory structures for providing various services such as scheduling, virtual memory, device IO and others. While this is necessary for enabling multi-user interactive systems (such as terminals), running specialized applications such as networking services do not require these features. Further, applications running on hypervisors face a duplication in abstraction layers, because the same abstractions are provided at both OS and hypervisor layer (though this can be offset with hardware acceleration such as KVM and SRIOV). But, the most important problem with an OS in the context of intra-application cache partitioning is that it renders any analysis of the application useless. When an application switches into a cache partition, it maybe preempted by the OS at any point of time. Now, the OS needs to switch its cache partition but it may still interfere with the memory buffers of the application (through cache flushes). In this case, even an effective intra-application partitioning scheme would not be useful in the presence of an OS. Our insight is that this overhead is mitigated by using unikernels, which offset the interference caused by the OS, since they are specialized kernels which bundle only the necessary platform code to run the application.

We provide an overview of our solution architecture in Fig.2. The boxes in green represent our own contribution and the grey boxes represent the existing components that we have reused. We start with the description of the set_partition application programming interface (API). The application has access to an intrinsic function called set_partition(int partition_id). This allows it to denote an abstract partition for each section of the program. We expect the programmer to instrument their applications with calls to this function using different partition IDs to denote which section the program is executing in. For example, a programmer




Figure 2: Solution architecture

may denote the networking parts of their application to execute in partition 1 and computational parts of their application execute in partition 2. Since these identifiers are abstract, the application can decide the scheme it uses to partition the application. A simple partitioning scheme would be to demarcate all network IO into a certain partition and the rest of the application into another. Further refinements could even create partitions within the networked section for more fine-grained control. Internally, this intrinsic function is implemented as a hypercall in the hypervisor. A hypercall is a mechanism for the user program to signal to the hypervisor that it needs access to a special resource. This hands over control to the hypervisor which performs the necessary operations to hand over control of the resource and returns to the user program. Next, we describe the unikernel library - Unikraft, that we use to turn our programs into unikernels. A unikernel library is a specialized library that provides the necessary platform code to turn an application in a bare-metal kernel. This kernel can now directly execute on hardware (though it requires a hypervisor to provide some basic platform abstractions). The advantage of this approach is that only the necessary code needed to run the application is pulled into the final application binary. Code for unnecessary features such as virtual memory, scheduling, IO emulation and others is not present in the kernel and so the memory footprint of these features is avoided. Unikraft [5] is a library OS, that provides the necessary platform code and a build system to turn regular applications into unikernels. These unikernels can be executed on a variety of hypervisors such as KVM and Xen. It has implementations of standard POSIX compatible system calls and external libraries can be ported to the Unikraft environment for additional features. All the applications that we wrote for performance evaluation were turned into unikernels using Unikraft and then executed on a hypervisor. We also had to port additional libraries to support our applications. Next, we provide details of the KVM hypervisor and how we integrated our cache partitioning scheme into the hypervisor. There are two types of hypervisors, Type-I and Type-II hypervisors. Type-I hypervisors run directly on the CPU and provide isolation and platform abstraction features. Type-II hypervisors are integrated with an existing operating system and reuse the features of the operating system to provide platform abstraction. Both types of hypervisors use hardware virtualization to provide isolation and fast emulation of platform features. In our setup, we have used the KVM hypervisor, which is integrated with the Linux operating system and uses Intel VT-d to provide hardware virtualization. KVM has support for intercepting hypercalls from guest kernels that it is running.

It is in this subsystem that we implement our cache-partitioning feature. We implement a new hypercall which intercepts calls to `set_partition` from the guest OS. Then, we inspect the partition ID provided by the guest OS and use it to allot it to a cache partition. An important nuance to keep in mind is that the guest OS may be executing on a different hardware thread than where the hypervisor intercepted it. Thus it is important to use the features of KVM to signal the correct core to allot the cache partition. Further, the choice of allocation policy also plays an important role in the effectiveness of the solution. Here we implement a simple round-robin scheme which allocates each new call to a different partition.




Figure 3: Example of cache bitmasks

Finally, this allotment is performed using the Intel Resource Director Technology (RDT). Intel RDT has a subfeature called Cache Allocation Technology (CAT) which is a hardware feature that allows application level cache partitioning. The first key component of CAT is the cache bitmask. A cache bitmask is a fixed size bitstring that is used to represent the regions of the cache that the application may occupy. We show some examples of bitmasks in Figure 3. As can be seen, a 1 represents a cache region that an application may occupy whereas a 0 represents a region it cannot. Thus, Bitmask 1 represents access to the full cache and Bitmasks 2 and 3 to smaller sections of it. Notice that Bitmask 2 and 3 are non-overlapping, thus applications execute under these bitmasks will not interfere with the cache footprint of each other. Each processor supports a fixed number of bitmasks to be preprogrammed in the system. Once this is done, we can write a model-specific register in each logical core which denotes which bitmask the current application should run under. This has the effect of partitioning the cache such that the requests from the core now go to the specific cache partition. We use this feature to dynamically switch cache partitions whenever an application requests a `set_partition` call.

4 EVALUATION AND RESULTS

The evaluation of our implementation requires careful analysis because there are two competing effects at play. On one hand, partitioning cache space for the networking parts may improve the effectiveness of DDIO by preventing the eviction of RX/TX buffers. But on the other hand, this may reduce the cache space available for the remaining parts of the application. Thus, even though the performance of the networking parts improves the overall application performance may suffer. Moreover, since our dynamic cache partitioning scheme is implemented using hypercalls in an hypervisor, there is an overhead associated with the VM exit. If these events are too frequent, it may add a non-negligible overhead leading to an increase in runtime. Given this, it would not be possible to determine the impact of our work by simply measuring wall-clock time.

While an improvement in wall-clock time is beneficial, it does not provide any insight into which part of our solution resulted in this improvement. In the case that performance is reduced, we need a method of diagnosing what led to the decrease in performance. To deal with this, we use statistical profiling of hardware performance counter events to benchmark different aspects of an application. We briefly hardware counter events and how they are recorded to better understand how we use it in our evaluation method. Modern processors have hardware counters, that count the number of occurrences of certain hardware events. These counters can capture counts ranging from CPU events such as instructions executed, cycles spent stalling and others, all the way upto system events such as memory requests, cache hits and misses and so on. Since constantly reading these counters will have a very high overhead, the processor can be programmed to generate an interrupt when a certain count is reached. Thus, this is referred to as statistical profiling, since we only capture a subset of the events. But the relative values of these counts can give us a good idea if the frequency of a certain event increased or decreased. Thus we use statistical profiling, through tools like `perf` to capture the frequency of certain events that we believe helps us understand the behavior of our application. First, we focus on PCIE Writes / Reads from the last level cache (LLC). This is the primary event that identifies if our solution is working. This event provides information about the number of PCIE accesses that were serviced from the LLC. This is important because if we are servicing more requests from the LLC, this implies that more network packets are staying resident in the cache and directly being processed from there, instead of being retrieved from main memory. Second, we focus on the total LLC Misses in the application. This is important because by partitioning the cache there is a chance we increase the number of cache misses in the application. Thus, we would like to analyze if the total number of application cache misses is increasing substantially. Finally, we also measure application level benchmarks such as bandwidth and latency of requests served, since finally any improvement in the software should result in an observed improvement in the client side performance. We have implemented four benchmarks - (1) **CS**: This program implements the Websocket protocol over TCP and creates a chat server. (2) **ML**: This program implements different ML algorithms to perform packet based quality-of-service classification (3) **KV**: This program implements a replicated key-value store using leader-server replication (4) **DL**: This program implements a deep learning inference workload for image recognition We believe this set of benchmarks are representative because they capture different types of workloads. **CS** is mostly network IO with little computation since it only performs Websocket packet-framing and forwarding over TCP. **ML** and **DL** represent a balance between network IO and compute, since the inference procedure requires non-trivial amounts of computation. Finally, **KV** represents a combination of network IO and disk IO, since the records are stored on the filesystem. Thus, evaluating the performance of our solution over these benchmarks should help us understand how it works. All of our benchmarks are executed on an AWS m5zn.metal machine with 48 logical cores, running on Amazon Linux 2, using QEMU 6.0.0 and Unikraft 0.10.0. We have 4 test scenarios - (1) **OS**: run the application on the OS, (2) **OS+PT**: run the application on the OS stack + fine-grained partitioning, (3) **UK**: run the application as a

unikernel, (4) **UK+PT**: run the application as a unikernel + fine-grained partitioning. This will help to evaluate the effect of each component of our solution. In the following figures, we present and discuss our results. Note that the absolute numbers of the hardware counters are normalized, to make comparison easier.

First, we observe the PCIE hits for each benchmark in Fig.4. Across all setups, we note that the PCIE hits increases. In the case of **OS+PT** case, this is not significant, implying that the OS interference is removing any benefits that we may observe. In the **UK** and **UK+PT** case, we observe that the PCIE hits increases. This shows that our cache partitioning has the effect of isolating network buffers, so that PCIE hits are increased. The rate of increase across different applications is different because of various network packet parameters such as size of packets, number of packets received and so on.




Figure 4: PCIE Hits for each benchmark

Next we move to the LLC misses observed for each benchmark. As shown in Fig. 5., LLC misses drop consistently when we move from **OS** to **UK**. This is because the smaller memory footprint of the unikernel allows more cache space for the application, thus reducing the number of cache misses. But, our partitioning scheme increases the cache misses, whether it is added to **OS** or **UK**. This is expected, since the computational parts of the application are receiving less cache space, more cache misses are expected. In the case of the compute heavy benchmarks such as **ML** and **DL**, we observe that cache misses are significantly increased, since they perform lot of computation. In the case of **KV** and **ML**, this is less since they are not performing too much computation and are mostly IO bound.




Figure 5: LLC Misses for each benchmark

In Fig. 6, we discuss the results of our throughput experiments. For each application, we setup multiple clients who then send requests to saturate the link and we check how many requests per second we can serve. We note that throughput increases for most benchmarks when going from **OS** to **UK** due to the smaller memory footprint and less cache contention. In the case of the **CS** and **KV** benchmarks we see an increase in the performance due to our partitioning scheme. But in the case of **DL** and **ML**, there is a drop in the performance. Our hypothesis is that since **CS**, **KV** are IO dominated, our scheme isolates the network buffer and improves the networking performance which translates to an improvement in the overall application performance. obtained. In the **DL** and **ML** case, it is possible that the smaller cache allocation is increasing the time taken to perform the computationally heavy part so much that any improvements in the networking part are mitigated. This is reasonable, since these programs are very sensitive to the cache size.




Figure 6: Throughput for each benchmark

Finally, we observe the P99 latency for the various benchmarks in Fig. 7. As observed, the P99 latency improves (smaller is better) for only the **CS** benchmark and gets worse for all others. A similar hypothesis from the throughput section applies, where any improvement in the networking parts of the code are mitigated by the rest of the application. Moreover, limiting the cache space may have the unforeseen effect of evicting some packets, which were previously able to be served from the cache, thus increasing the P99 latency.




Figure 7: P99 Latency for each benchmark

5 RELATED WORK

The work most closely related to our work is [2]. The work performs a thorough analysis of the Intel DDIO implementation. Their experiments consider different types of applications executed under different cache partitioning schemes. The major takeaway from their work that we capitalized on was that there is no *one-size-fits-all* cache partitioning policy. Thus we implement a solution that allows an application to customize its use of the cache, tailored to its specific needs. On the front of fine-grained cache-partitioning, there are two main directions - software based cache partitioning based on colored pages [4], [6] and hardware based way partitioning (such as Intel RDT) [1], [3]. The work closest to our implementation is [8], which also proposes a fine-grained cache partitioning scheme. But their work does not consider the effect of this scheme on networked applications. Further the implementation is not performed using Intel RDT (which is a more mainstream implementation of hardware cache partitioning) and also does not take into the effect of hypervisor (necessary for safe isolation of applications).

6 CONCLUSION AND FUTURE WORK

We have designed and implemented a solution to reduce the effect of cache contention on the performance of network IO and demonstrated our results. While the results are not promising, since we do not see a consistent improvement in performance across all the benchmarks, the **CS** and **KV** benchmarks provide insight into situations where this solution may work. Additionally, it also corroborates the claim of [2] that heavily CPU or memory bound applications do not benefit much from DDIO tuning based on cache partitioning as observed in **ML** and **DL**. Still, I feel this work opens up interesting future directions. The first and most important learning is that this technique can be specialized for specific applications such as virtual switches, network function virtualization and similar applications which are dominated by network IO. Further, the cache allocation policy needs a thorough review to ensure that it is performing optimally. This would require digging deeper into the cache allocation being performed by our simple round-robin algorithm and how it can be improved. On the measurement front, we feel there are multiple things that can be improved to make this work better. First, the network stack used in the **OS** and **UK** setup is different, since we could only port a simple TCP/IP stack to the unikernel environment. It would be good to have something like DPDK tuned for both the setups, thus giving a better comparison. If a fair baseline is established between the **OS** and unikernel setup, then we can be confident that our implementation improves the performance of the application. Second, the measurements are performed in noisy environments, because the Linux kernel is still running on the cores that are running the unikernels. Isolating these cores requires the use of advanced features such as ISOLC-PU, which are not available in the EC2 instances that we used. Thus, we believe that we have made a case for intra-application cache partitioning and unikernels.

ACKNOWLEDGMENTS

To Prof. Muhammad Shahbaz for conducting such a great course and Sharuna Anandraj, Darwin Kamanuru, K M A Solaiman and Kai Ling for helping out throughout the semester!

REFERENCES

- [1] Ruobing Chen, Jinping Wu, Haosen Shi, Yusen Li, Haiyan Yin, Shanjian Tang, Xiaoguang Liu, and Gang Wang. 2020. Deep Learning Assisted Resource Partitioning for Improving Performance on Commodity Servers. In *Proceedings of the ACM International Conference on Parallel Architectures and Compilation Techniques* (Virtual Event, GA, USA) (PACT '20). Association for Computing Machinery, New York, NY, USA, 153–154. <https://doi.org/10.1145/3410463.3414668>
- [2] Alireza Farshin, Amir Roozbeh, Gerald Q. Maguire, and Dejan Kostić. 2020. Reexamining Direct Cache Access to Optimize I/O Intensive Applications for Multi-Hundred-Gigabit Networks. In *Proceedings of the 2020 USENIX Conference on Usenix Annual Technical Conference (USENIX ATC'20)*. USENIX Association, USA, Article 46, 17 pages.
- [3] Liran Funaro, Orna Agmon Ben-Yehuda, and Assaf Schuster. 2016. Ginseng: Market-Driven LLC Allocation. In *Proceedings of the 2016 USENIX Conference on Usenix Annual Technical Conference (Denver, CO, USA) (USENIX ATC '16)*. USENIX Association, USA, 295–308.
- [4] Jorg Herter, Peter Backes, Florian Hauptenthal, and Jan Reineke. 2011. CAMA: A Predictable Cache-Aware Memory Allocator. In *Proceedings of the 2011 23rd Euromicro Conference on Real-Time Systems (ECRTS '11)*. IEEE Computer Society, USA, 23–32. <https://doi.org/10.1109/ECRTS.2011.11>
- [5] Simon Kuenzer, Vlad-Andrei Bădoiu, Hugo Lefeuvre, Sharan Santhanam, Alexander Jung, Gauthier Gain, Cyril Soldani, Costin Lupu, Ștefan Teodorescu, Costi Răducanu, Cristian Banu, Laurent Mathy, Răzvan Deaconescu, Costin Raiciu, and Felipe Huici. 2021. Unikraft: Fast, Specialized Unikernels the Easy Way. In *Proceedings of the Sixteenth European Conference on Computer Systems* (Online Event, United Kingdom) (EuroSys '21). Association for Computing Machinery, New York, NY, USA, 376–394. <https://doi.org/10.1145/3447786.3456248>
- [6] Timothy Sherwood, Brad Calder, and Joel Emer. 1999. Reducing Cache Misses Using Hardware and Software Page Placement. In *Proceedings of the 13th International Conference on Supercomputing* (Rhodes, Greece) (ICS '99). Association for Computing Machinery, New York, NY, USA, 155–164. <https://doi.org/10.1145/305138.305189>
- [7] Amin Tootoonchian, Aurojit Panda, Chang Lan, Melvin Walls, Katerina Argyraki, Sylvia Ratnasamy, and Scott Shenker. 2018. ResQ: Enabling SLOs in Network Function Virtualization. In *Proceedings of the 15th USENIX Conference on Networked Systems Design and Implementation* (Renton, WA, USA) (NSDI'18). USENIX Association, USA, 283–297.
- [8] Xiaodong Wang, Shuang Chen, Jeff Setter, and José F. Martínez. 2017. SWAP: Effective Fine-Grain Management of Shared Last-Level Caches with Minimum Hardware Support. In *2017 IEEE International Symposium on High Performance Computer Architecture (HPCA)*. 121–132. <https://doi.org/10.1109/HPCA.2017.65>

A CONTRIBUTIONS

A.1 Jiangqiong Liu

Worked on the handling the measurement across the entire project. Implemented perf scripts to automate collection of hardware counter events for all applications and perform the post-processing to report the results. Correlated performance changes in application based on measurement results.

A.2 Joshua Repogle

Implemented an end-to-end Websocket server in C++ and used it to implement a chat server. The C++ server was written without dynamic memory allocation to carefully control the placement of memory buffers. Identified networking parts of the application and performed `set_partition` instrumentation.

A.3 Vedaant Rajoo

Implemented a replicated, in-memory key-value store using C++. Analyzed networking and computational parts of the application and performed `set_partition` instrumentation.

A.4 Sowmya Jayaram Iyer

Implemented machine learning models for packet based classification. Executed these models on real-world traces and ported them

to run as Unikraft kernels. Identified the `set_partition` instrumentation points for compute bound parts of the application.

A.5 Keerthana Ashokkumar

Designed and implemented a replicated, in-memory key-value store using Java. Implemented two phase commit based atomic updates and leader server replication. Assisted in the C++ implementation of the in-memory store and identified networking parts of the application.

A.6 Anmol Sahoo

Designed and implemented the Intel RDT based cache partitioning scheme as a KVM patch in the Linux kernel. Exposed the userspace API as VMCALLs in KVM. Ported the necessary libraries to Unikraft (rplib for rpc, tcp/ip for sockets, tvml for deep learning / machine learning). Performed end-to-end execution and testing of all applications.